



INTRODUCTION TO

Neural Networks in Python

Sabriya Alam and Parker Crain
Purdue University
Dec. 19th 2019

Outline

- Neural networks overview
- Data preprocessing techniques
- Base neural network model
- Network tuning
- Model metrics and evaluation



Neural Network Overview

<https://ujiwalkarn.me/2016/08/09/quick-intro-neural-networks/>




- What is an artificial neural network?
 - Machine learning computational model
 - Inspired by how biological neural networks in our brains process information and draw conclusions.
- Neural networks are used for
 - Speech recognition
 - Computer vision (image / video processing)
 - Text processing
 - Data analytics
- Many types of neural networks exist
 - Our focus will be on multi-layer perceptron



Neural Network Overview


<https://engineering.purdue.edu/~milind/ece20875/2019fall/notes/lecture-22.pdf>

- Goal is to be able to learn non-linear classification boundaries
- Fundamental building block is a neuron (or perceptron)
 - Node with inputs and weights
 - Dot product is input to activation function
 - Output of neuron is result of activation function
- Single perceptron will converge if linear decision boundary exists. Multi-layer required for non-linear.
- Multi-layer = outputs of neurons become inputs to another layer of neurons
- Error in predicted values are back-propagated to adjust weights in the correct direction



Data Preprocessing Techniques

- Data comes in two types:
 - Labelled
 - Used for supervised machine learning
 - Known outcomes
 - Unlabelled
 - Used for unsupervised (deep) learning
 - Finds patterns without knowing how many classes there are or what class each data point belongs to
 - Examples: audio recordings, image data sets, etc.
- Our data set is labelled



Data Preprocessing : Balancing Data

<https://datascience.aero/predicting-improbable-part-1-imbalanced-data-problem/>

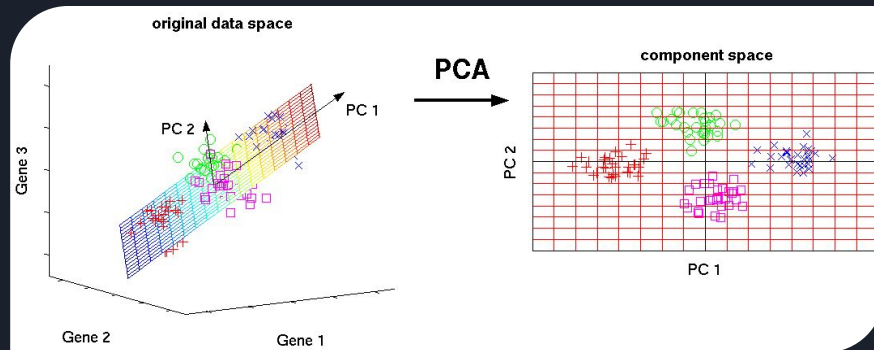
- Data needs to be divided into training and testing sets
- Eliminate bias (skewed results) by creating evenly balanced sets
 - If training set is 70% class A and 30% class B, the model may be biased (at least initially) toward predicting class A
 - If you know that this proportion is how the real world works, and the probability of event A occurring truly is 70%, then this is a relevant factor that should be kept in
 - Otherwise, create balanced sets with 50% class A and 50% class B



Data Preprocessing : Dimensionality Reduction

<https://towardsdatascience.com/pca-using-python-scikit-learn-e653f8989e60>

- Having a lot of input factors can slow down a model
- To make it faster, eliminate “duplicate” variables
 - Linearly correlated variables do not give any additional information, so keep only one
- Many statistical techniques for this
 - lasso regression
 - principal component analysis
- Using PCA brought our relevant features from 89 to 60



Code : PCA of balanced data sets

```
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import pandas as pd
import numpy as np
```

Import packages for
normalizing data and
performing PCA

```
def preprocess(filename, filename2):
    dataset = pd.read_csv(filename, skiprows=2)
    training_array = np.array(dataset)
    X_train = training_array[:, 1:90]
```

Import training data
from *filename* as
shown in previous
presentation

Import testing data
from *filename2*

PCA requires
normalized data, so
use scalar to make
adjustments

```
    testset = pd.read_csv(filename2)
    test_array = np.array(testset)
    X_test = test_array[:, 1:90]

    scaler = StandardScaler()
    scaler.fit(X_train)
    X_train = scaler.transform(X_train)
    X_test = scaler.transform(X_test)
```

```
    pca = PCA(.95)
    pca.fit(X_train)
    n = pca.n_components_
```

Create PCA model that will
minimize number of
parameters while still
preserving 95% of the
variance in the data

```
    train_pca = pca.transform(X_train)
    test_pca = pca.transform(X_test)
```

```
    return train_pca, test_pca, n
```

Apply PCA
transformation to the
two data sets

Neural Network Libraries in Python

<https://towardsdatascience.com/introduction-to-deep-learning-with-keras-17c09e4f0eb2>

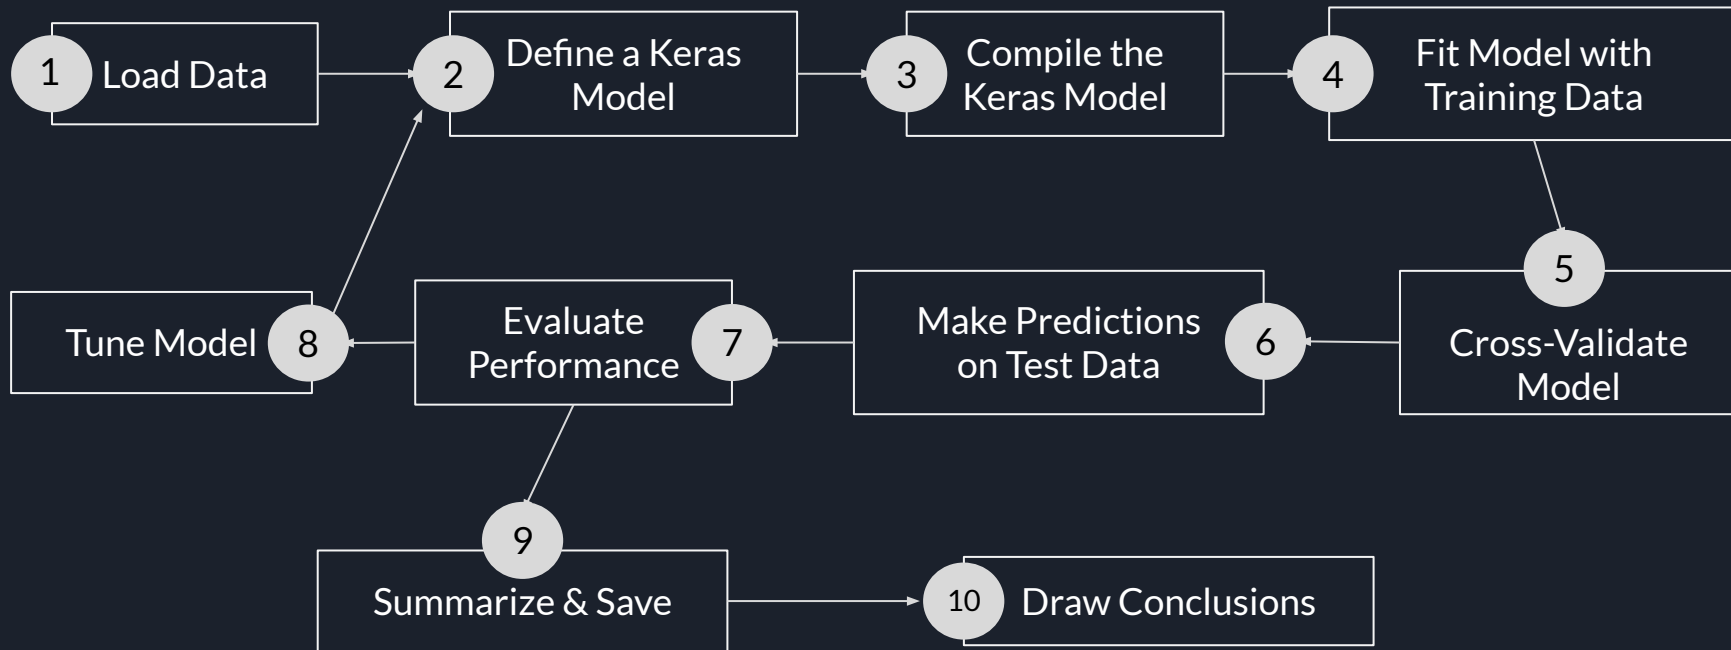
<https://medium.com/@Vatsal410/keras-without-nvidia-gpus-with-plaidml-and-amd-gpu-4ba6f60025ce>



- Many options for neural network libraries in Python
- We chose to use Keras
 - Open-source
 - Can run on top of TensorFlow, R, PlaidML, Microsoft Cognitive Toolkit, Theano, etc.
 - User friendly
- Default backend is TensorFlow
- Without an Nvidia GPU, it is better to use PlaidML for accelerated performance


General Workflow

<https://machinelearningmastery.com/tutorial-first-neural-network-python-keras/>



Base Neural Network Model

<https://machinelearningmastery.com/tutorial-first-neural-network-python-keras/>



- There are no rules of thumb for choosing neural network model parameters.
 - Unique to each data set
 - Depends on data scope, structure, diversity and complexity
- Start with some randomized values, and then tune the model based on its performance.
- Base neural network is usually an input layer, one hidden layer, and one output layer.

Base Neural Network Model

Set up PlaidML and Keras, and import necessary modeling tools

Read in training and testing data, along with expected outputs

3 Layers :
Input, Hidden, and Output

As the model is being built, the accuracy is being cross validated with subsections of the training data too

```
import plaidml.keras
plaidml.keras.install_backend()
from keras.models import Sequential
from keras.layers import Dense
```

```
def main(filename, filename2):
    n = 89
    dataset = pd.read_csv(filename, skiprows=2)
    new_array = np.array(dataset)
    X = new_array[:,1:90]
    y = new_array[:, 90]
    dataset = pd.read_csv(filename2)
    new_array2 = np.array(dataset)
    X_test = new_array2[:, 1:90]
    y_test = new_array2[:, 90]
```

```
model = Sequential()
model.add(Dense(10, input_dim=n, activation='relu'))
model.add(Dense(8, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(X, y, validation_split=0.33, epochs=5, batch_size=10)
_, accuracy = model.evaluate(X, y)
print('Accuracy: %.2f' % (accuracy * 100))
```

Terminology:

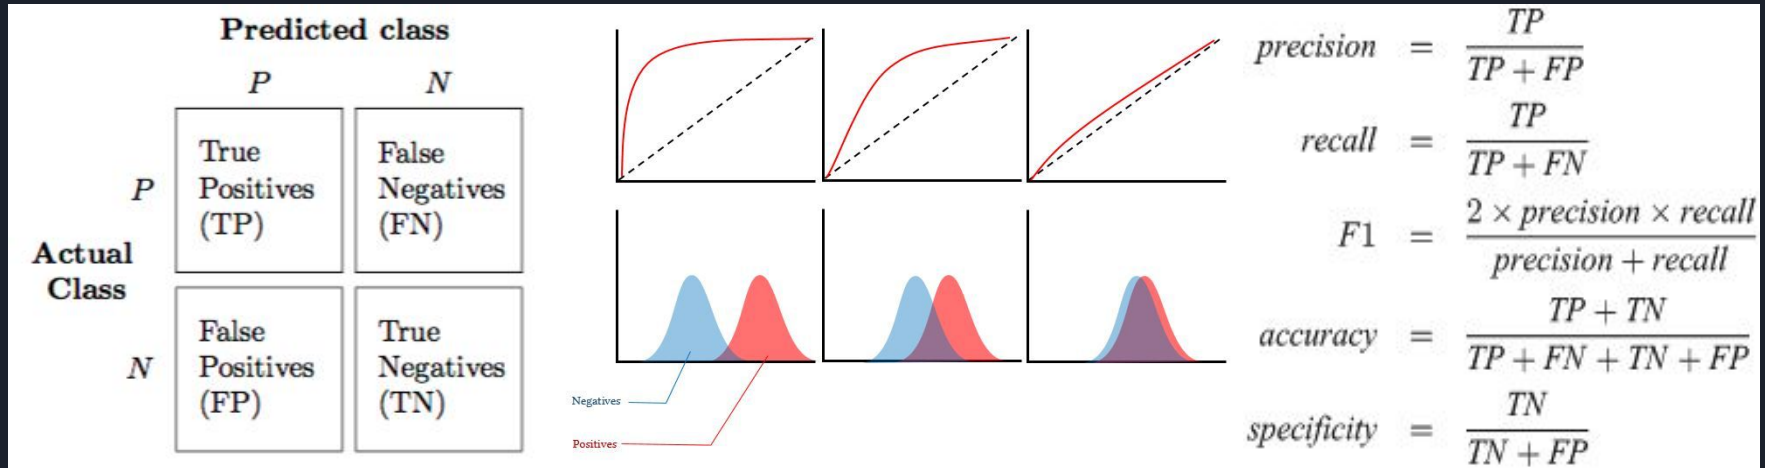
Sequential - model is linear stack of layers

Dense - densely connected NN layer, so every node of this layer will connect to every node of the next layer

Model Performance Metrics & Evaluation

<https://androidkt.com/get-the-roc-and-auc-for-keras-model/>

- Some metrics used for evaluating neural networks
 - ROC curve - plots false positive (x) and true positive (y) rates at various thresholds
 - Diagonal is equal to random guessing
 - Area under curve approaches 1 for perfect classifier
 - Confusion matrix
 - Precision, recall, and F1 scores
- Keras has built in functions to generate these



Basic Model : Evaluation Metrics

```
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_curve, roc_auc_score
from keras.models import load_model
```

Can view node weights

```
print(model.get_weights())
print(type(model.get_weights()))
print("Model Summary:")
print(model.summary())
scores = model.evaluate(X, y, verbose=0)
print("%s: %.2f%%" % (model.metrics_names[1],
scores[1] * 100))
```

Model summary shows number of layers and nodes

Save model, so you do not need to regenerate it every time

```
# save model and architecture to file
model.save("modelbasic.h5")
print("Saved model to disk")

print("loading model")
# load model
model = load_model('modelbasic.h5')
```

Input test data and get predictions from the model.

Print for comparison.

Calculate false positive and true positive rates for plotting ROC and getting AUC score

Get precision, recall, f1-score, and confusion matrix.

```
import matplotlib.pyplot as plt


def plot_roc_curve(fpr, tpr):
    plt.plot(fpr, tpr)
    plt.axis([0, 1, 0, 1])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.show()
```

```
# make class predictions with the model
pred = model.predict_classes(X_test)
for i in range(0, len(pred)):
    print('%d (expected %d)' % (pred[i],
y_test[i]))

fpr, tpr, thresholds = roc_curve(y_test, pred)
plot_roc_curve(fpr, tpr)
print("auc score")
auc_score = roc_auc_score(y_test, pred)
print(auc_score)

print(classification_report(y_test, pred))
print("confusion matrix:")
print(confusion_matrix(y_test, pred))
```

Basic Model Evaluation Statistics



Accuracy: 60.02%

AUC Score = 0.5853871128871129

Classification Summary

	precision	recall	f1-score	support
0.0	0.55	0.94	0.69	1000
1.0	0.80	0.23	0.35	1001
accuracy			0.59	2001
macro avg	0.68	0.59	0.52	2001
weighted avg	0.68	0.59	0.52	2001

Confusion Matrix:

[[974 26]

[835 166]]

Interpretation

Accuracy: 60.02%
AUC Score = 0.5853871128871129

Classification Summary


	precision	recall	f1-score	support
0.0	0.55	0.94	0.69	1000
1.0	0.80	0.23	0.35	1001
accuracy			0.59	2001
macro avg	0.68	0.59	0.52	2001
weighted avg	0.68	0.59	0.52	2001

Confusion Matrix:

[[974 26]
[835 166]]

- ROC curve is not far off from the diagonal, indicating that the classifier is not much better than randomly guessing.
- Area under curve is .58538, but we want this number to be as close to 1 as possible.
- We need to tune the model!


- If the model says the patient will not be readmitted, probability that they will actually not be readmitted is .55
- If the model says the patient will be readmitted, the probability that they will actually be readmitted is .80
- If the patient is not readmitted, the probability of the model telling me so is .94
- If the patient is readmitted, the probability of the model telling me so is .23
- Confusion matrix indicates that of the 2001 data points used for testing, there were 974 true negatives, 166 true positives, 26 false positives, and 835 false negatives.



Network Tuning

<https://machinelearningmastery.com/how-to-configure-the-number-of-layers-and-nodes-in-a-neural-network/>

- What factors improve a model?
- Deeper networks tend to perform better
 - Increase number of layers
 - Adjust number of nodes per layer
 - Try various activation functions
- Adjust number of epochs
 - Entire training set passed through network
 - Weights all updated once per epoch
 - Be careful about overfitting to training set
- Iterative experimentation process



Tuned Model

```
model = Sequential()
model.add(Dense(30, input_dim=n, activation='relu'))
model.add(Dense(20, activation='relu'))
model.add(Dense(20, activation='relu'))
model.add(Dense(15, activation='relu'))
model.add(Dense(15, activation='relu'))
model.add(Dense(15, activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(15, activation='relu'))
model.add(Dense(8, activation='sigmoid'))
model.add(Dense(5, activation='sigmoid'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(X, y, epochs=200, batch_size=10)
_, accuracy = model.evaluate(X, y)
print('Accuracy: %.2f' % (accuracy * 100))
```

Increased number of layers
and nodes

Variation in activation
functions

Epochs increased to 200

Tuned Model : Evaluation Metrics

```
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_curve, roc_auc_score
from keras.models import load_model
```

Can view node weights

```
print(model.get_weights())
print(type(model.get_weights()))
print("Model Summary:")
print(model.summary())
scores = model.evaluate(X, y, verbose=0)
print("%s: %.2f%%" % (model.metrics_names[1],
scores[1] * 100))
```

Model summary shows number of layers and nodes

Save model, so you do not need to regenerate it every time

```
# save model and architecture to file
model.save("modeltuned.h5")
print("Saved model to disk")

print("loading model")
# load model
model = load_model('modeltuned.h5')
```

Input test data and get predictions from the model.

Print for comparison.

Calculate false positive and true positive rates for plotting ROC and getting AUC score

Get precision, recall, f1-score, and confusion matrix.

```
import matplotlib.pyplot as plt


def plot_roc_curve(fpr, tpr):
    plt.plot(fpr, tpr)
    plt.axis([0, 1, 0, 1])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.show()
```

```
# make class predictions with the model
pred = model.predict_classes(X_test)
for i in range(0, len(pred)):
    print('%d (expected %d)' % (pred[i],
y_test[i]))

fpr, tpr, thresholds = roc_curve(y_test, pred)
plot_roc_curve(fpr, tpr)
print("auc score")
auc_score = roc_auc_score(y_test, pred)
print(auc_score)

print(classification_report(y_test, pred))
print("confusion matrix:")
print(confusion_matrix(y_test, pred))
```

Tuned Model Evaluation Statistics



Accuracy: 67.28%

AUC score = 0.6601883116883117

Classification Report

	precision	recall	f1-score	support
0.0	0.65	0.70	0.67	1000
1.0	0.67	0.62	0.65	1001

accuracy		0.66	2001	
macro avg	0.66	0.66	0.66	2001
weighted avg	0.66	0.66	0.66	2001

Confusion Matrix:

[[697 303]

[377 624]]

Interpretation

Accuracy: 67.28%
AUC score = 0.6601883116883117

Classification Report

	precision	recall	f1-score	support
0.0	0.65	0.70	0.67	1000
1.0	0.67	0.62	0.65	1001
<hr/>				
accuracy		0.66		2001
macro avg	0.66	0.66	0.66	2001
weighted avg	0.66	0.66	0.66	2001


Confusion Matrix:

[[697 303]

[377 624]]


- ROC curve is farther from the diagonal, indicating that the classifier has more reliable predictive power.
- Area under curve is .66018, which is closer to 1 than before

- If the model says the patient will not be readmitted, probability that they will actually not be readmitted is .65
- If the model says the patient will be readmitted, the probability that they will actually be readmitted is .67
- If the patient is not readmitted, the probability of the model telling me so is .70
- If the patient is readmitted, the probability of the model telling me so is .62
- Confusion matrix indicates that of the 2001 data points used for testing, there were 697 true negatives, 624 true positives, 303 false positives, and 377 false negatives.




Comparison of Base and Tuned Models

- ROC shows second model is better at discerning between positive and negative classes
 - Area increased from .58 to .66
- Accuracy of predictions on the same testing data set increased from 60.02% to 67.28%.
- Proportion of true positive and true negative predictions are about the same in second model
 - Initial model had skewed results (especially inaccurate in predicting positive cases).
- F1 score combines precision and recall (harmonic mean), and a model should maximize this to have an optimal balance of these two metrics. Should be as close to 1 as possible.
 - Original model f1 scores: .69 for classification 0 & .35 for classification 1
 - Improved model f1 scores: .67 for classification 0 & .65 for classification 1.
- F1 score should be close to 1 for both classes to indicate a better classifier--confirms second model as more reliable predictive model.



Can we predict a patient's likelihood of hospital readmission?



- This model indicates that there may be a correlation between a patient's lifestyle and chances of being readmitted.
- Accuracy is not high enough to be definitive, but a model may be a helpful tool in tandem with a professional's advice.
- However, neural network is more accurate in this case than regression models were, which is promising.
- With more exhaustive tuning of this model, there is a prospect of creating a model that is increasingly reliable in providing predictions.

Concluding Thoughts

